

CGAN Project


Project Presentation

Presented by: *Mabin Varghese*


Task And Objective

- Dataset: Flying objects -> set of frames
- Extract FF and LF of every sets (Total 300 sets)
- Build a Generative adversarial network

Training input/Real images

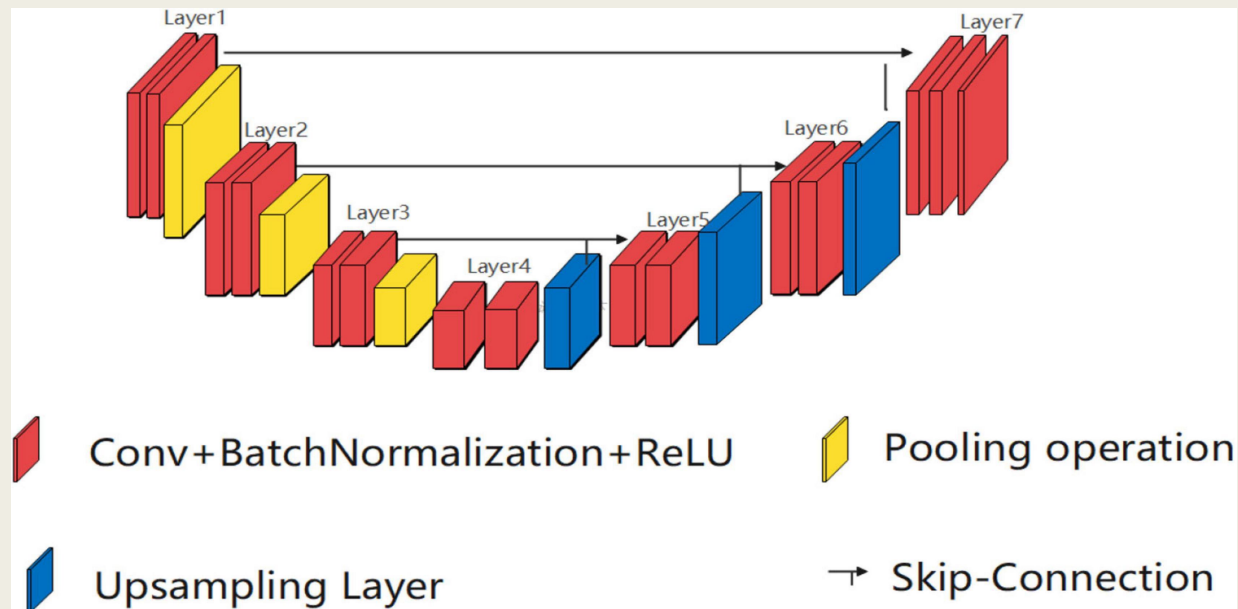


Last Frames



Model Architecture

1. Generator Model
 - i. UNET generator model -> Typical for Image segmentation
 - ii. Encoder and decoder
 - iii. Skip connections -> preserve spatial features and capture low- and high-level feature
2. Noise – Random noise with a normal distribution

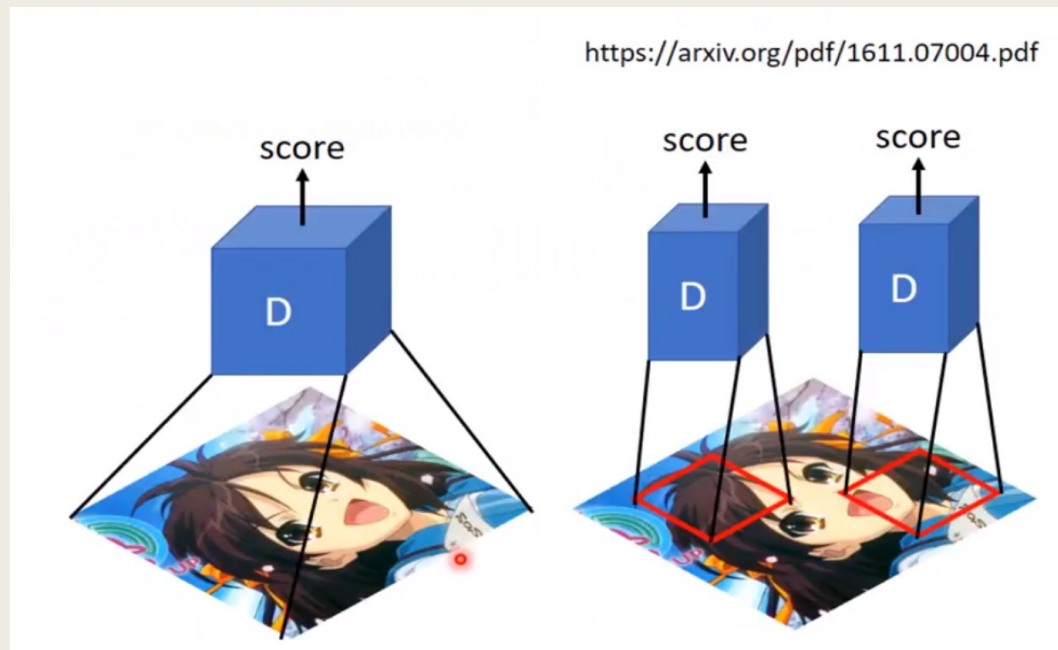


https://www.frontiersin.org/files/Articles/841297/fnagi-14-841297-HTML-r2/image_m/fnagi-14-841297-g001.jpg

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 32, 32, 32]	5,792
Conv2d-2	[-1, 64, 32, 32]	18,496
Conv2d-3	[-1, 128, 16, 16]	137,344
BatchNorm2d-4	[-1, 128, 16, 16]	256
ReLU-5	[-1, 128, 16, 16]	0
Conv2d-6	[-1, 256, 8, 8]	524,544
BatchNorm2d-7	[-1, 256, 8, 8]	512
ReLU-8	[-1, 256, 8, 8]	0
Conv2d-9	[-1, 512, 4, 4]	2,097,664
BatchNorm2d-10	[-1, 512, 4, 4]	1,024
ReLU-11	[-1, 512, 4, 4]	0
Conv2d-12	[-1, 512, 2, 2]	4,194,816
BatchNorm2d-13	[-1, 512, 2, 2]	1,024
ReLU-14	[-1, 512, 2, 2]	0
Conv2d-15	[-1, 512, 1, 1]	4,194,816
BatchNorm2d-16	[-1, 512, 1, 1]	1,024
ReLU-17	[-1, 512, 1, 1]	0
ConvTranspose2d-18	[-1, 512, 2, 2]	4,194,816
BatchNorm2d-19	[-1, 512, 2, 2]	1,024
ReLU-20	[-1, 512, 2, 2]	0
Dropout2d-21	[-1, 512, 2, 2]	0
ConvTranspose2d-22	[-1, 256, 4, 4]	4,194,560
BatchNorm2d-23	[-1, 256, 4, 4]	512
ReLU-24	[-1, 256, 4, 4]	0
Dropout2d-25	[-1, 256, 4, 4]	0
ConvTranspose2d-26	[-1, 128, 8, 8]	1,572,992
BatchNorm2d-27	[-1, 128, 8, 8]	256
ReLU-28	[-1, 128, 8, 8]	0
Dropout2d-29	[-1, 128, 8, 8]	0
ConvTranspose2d-30	[-1, 64, 16, 16]	393,280
BatchNorm2d-31	[-1, 64, 16, 16]	128
ReLU-32	[-1, 64, 16, 16]	0
Dropout2d-33	[-1, 64, 16, 16]	0
ConvTranspose2d-34	[-1, 32, 32, 32]	98,336
BatchNorm2d-35	[-1, 32, 32, 32]	64
ReLU-36	[-1, 32, 32, 32]	0
Dropout2d-37	[-1, 32, 32, 32]	0
ConvTranspose2d-38	[-1, 3, 32, 32]	867
BatchNorm2d-39	[-1, 3, 32, 32]	6
Tanh-40	[-1, 3, 32, 32]	0

Discriminator Model

- i. Patch GAN discriminator model -> maps a NxN array to the patch from the real images.
- ii. Series of convolutional layers, batch normalization, leaky ReLU activations, dropout layers, and a final classification layer with a sigmoid activation function.
- iii. Input the condition in discriminator also



```
discriminator = Disc()  
discriminator = discriminator.to(device)  
summary(discriminator, input_size=[(6,32,32)], device=device.type)
```

Layer (type)	Output Shape	Param #
Conv2d-1	[-1, 64, 16, 16]	6,208
BatchNorm2d-2	[-1, 64, 16, 16]	128
LeakyReLU-3	[-1, 64, 16, 16]	0
Conv2d-4	[-1, 128, 8, 8]	131,200
BatchNorm2d-5	[-1, 128, 8, 8]	256
LeakyReLU-6	[-1, 128, 8, 8]	0
Dropout2d-7	[-1, 128, 8, 8]	0
Conv2d-8	[-1, 256, 4, 4]	524,544
BatchNorm2d-9	[-1, 256, 4, 4]	512
LeakyReLU-10	[-1, 256, 4, 4]	0
ZeroPad2d-11	[-1, 256, 6, 6]	0
Dropout2d-12	[-1, 256, 6, 6]	0
Conv2d-13	[-1, 512, 3, 3]	2,097,664
BatchNorm2d-14	[-1, 512, 3, 3]	1,024
LeakyReLU-15	[-1, 512, 3, 3]	0
ZeroPad2d-16	[-1, 512, 5, 5]	0
Dropout2d-17	[-1, 512, 5, 5]	0
Conv2d-18	[-1, 1, 2, 2]	8,193
BatchNorm2d-19	[-1, 1, 2, 2]	2
AvgPool2d-20	[-1, 1, 1, 1]	0
Sigmoid-21	[-1, 1, 1, 1]	0

Total params: 2,769,731

Trainable params: 2,769,731

Non-trainable params: 0

Dataset & Transformation

Original Dataset

- Training – 300 pairs
- Validation – 60 pairs
- Testing – 60 pairs


Augmentation -> to increase pairs

- Color jitter
- Gaussian blur


New Dataset

- Training – 882
- Validation – 252
- Testing – 126


Training input/Real images



Last Frames



Generated images/Fake images



First Epoch Before Training


Conditional GAN Training

1. Initialized Xavier weights
2. Applied learning rate scheduling based on validation losses
3. Simultaneous Gen() and Disc() Training


Evaluation while training

1. Losses
2. L1 Loss – Assessing similarity b/w generated image and real image
3. SSIM scores
4. Visual inspection


Training input/Real images



Last Frames



Generated images/Fake images




Visualization on the 1st Epoch


SSIM Score : -0.083

After 10 Epoch



Training input/Real images



Last Frames




Generated images/Fake images




Epoch 6/500, Val Disc Loss: 6.6192686557769775, Val Gen Loss: 2.437118724698112
 Final Average L1 Loss: Test-: 299.7445445212107 Val-: 302.6081597994244 Train-: 301.9561856782355
 Epoch 7/500, Discriminator Loss: 6.603388356513717, Generator Loss: 5.729565869112944
 Epoch 7/500, Val Disc Loss: 6.618022209122068, Val Gen Loss: 2.384256725273435
 Final Average L1 Loss: Test-: 295.1511125715952 Val-: 297.9273597399394 Train-: 297.38670709181804
 Epoch 8/500, Discriminator Loss: 6.603120730307096, Generator Loss: 5.6475960073016935
 Epoch 8/500, Val Disc Loss: 6.640234069218712, Val Gen Loss: 2.287672980437203
 Final Average L1 Loss: Test-: 286.8669184427413 Val-: 289.99678199253384 Train-: 289.29934480022683
 Epoch 9/500, Discriminator Loss: 6.6028856771603195, Generator Loss: 5.567617116331243
 Epoch 9/500, Val Disc Loss: 6.641473514693123, Val Gen Loss: 2.2058519106062633
 Final Average L1 Loss: Test-: 281.4507105993846 Val-: 284.54173178899856 Train-: 283.7893042434641
 Epoch 10/500, Discriminator Loss: 6.60265444898281, Generator Loss: 5.4932345203261255
 Epoch 10/500, Val Disc Loss: 6.638700905300323, Val Gen Loss: 2.1360248563781616
 Final Average L1 Loss: Test-: 275.43047590861244 Val-: 278.2971281853933 Train-: 277.7097190048149
 Testing DT - Last 10 Average L1 Loss Trend: Decreasing
 Average SSIM Score for Training: 38.12792512018833
 Average SSIM Score for Validation: 38.02091590049586
 Average SSIM Score for testing: 38.626941468241036

After 70 Epoch


Training input/Real images




Last Frames



Generated images/Fake images




Generator Loss




Final Average L1 Loss: Test-: 115.05315701166788 Val-: 115.71146666057525 Train-: 110.62344776935318
 Epoch 68/500, Discriminator Loss: 6.601439960419185, Generator Loss: 3.6098242481819898
 Epoch 68/500, Val Disc Loss: 6.601313844559685, Val Gen Loss: 0.3593471448218066
 Final Average L1 Loss: Test-: 111.88131357942308 Val-: 112.47358329239347 Train-: 107.2018522591818
 Epoch 69/500, Discriminator Loss: 6.60143817633458, Generator Loss: 3.602707237343129
 Epoch 69/500, Val Disc Loss: 6.601471087289235, Val Gen Loss: 0.3460363337090091
 Final Average L1 Loss: Test-: 108.53355484349387 Val-: 109.2031237388414 Train-: 103.68656983721554
 Epoch 70/500, Discriminator Loss: 6.60143800629278, Generator Loss: 3.593542151440298
 Epoch 70/500, Val Disc Loss: 6.600648022833325, Val Gen Loss: 0.3390749206855183
 Final Average L1 Loss: Test-: 107.6504233337584 Val-: 108.5324189256108 Train-: 103.3987434340172
 Testing DT - Last 10 Average L1 Loss Trend: Not Strictly Decreasing
 Average SSIM Score for Training: 83.09536146627163
 Average SSIM Score for Validation: 80.83135344953132
 Average SSIM Score for testing: 81.05532626800729

After 450 Epoch


Training input/Real images




Last Frames



Generated images/Fake images




Generator Loss




Epoch 448/500, Discriminator Loss: 6.601401940494979, Generator Loss: 3.3305991398774575
 Epoch 448/500, Val Disc Loss: 6.601404954516699, Val Gen Loss: 0.12546464087559828
 Final Average L1 Loss: Test-: 16.603139894349233 Val-: 17.3287567431255 Train-: 11.252043261921324
 Epoch 449/500, Discriminator Loss: 6.601401899947601, Generator Loss: 3.3294837558620913
 Epoch 449/500, Val Disc Loss: 6.601388634197296, Val Gen Loss: 0.12607862775598372
 Final Average L1 Loss: Test-: 16.82409765346656 Val-: 17.555720793704193 Train-: 11.554057783787213
 Epoch 450/500, Discriminator Loss: 6.601401940494979, Generator Loss: 3.3293133960559524
 Epoch 450/500, Val Disc Loss: 6.601415219761077, Val Gen Loss: 0.12735636028948996
 Final Average L1 Loss: Test-: 16.99660283823808 Val-: 17.588565214758827 Train-: 11.501546374726999
 Testing DT - Last 10 Average L1 Loss Trend: Not Strictly Decreasing
 Average SSIM Score for Training: 94.72890512918323
 Average SSIM Score for Validation: 91.70754764283377
 Average SSIM Score for testing: 91.94685407911362

Evaluating Test Dataset


Training input/Real images



Last Frames




Generated images/Fake images




First Frame


Training input/Real images



Last Frames



Generated images/Fake images




Generated
Last Frame


Evaluating Test Dataset

```
Epoch 98/100, Discriminator Loss: 6.601402467610884, Generator Loss: 3.3399948862945146
Epoch 98/100, Val Disc Loss: 6.6013446875980915, Val Gen Loss: 0.07010654413274356
Final Average L1 Loss: Test-: 6.1906495737651035 Val-: 6.892448808583948 Train-: 3.394786059400257
Epoch 99/100, Discriminator Loss: 6.601402737926734, Generator Loss: 3.340074588922687
Epoch 99/100, Val Disc Loss: 6.601355851642669, Val Gen Loss: 0.06552046896623713
Final Average L1 Loss: Test-: 6.170129063465293 Val-: 6.630496521081243 Train-: 3.791787400799476
Epoch 00106: reducing learning rate of group 0 to 4.4372e-04.
Discriminator Learning Rate changed: 0.000499236317115515 -> 0.0004437212386522697
Epoch 100/100, Discriminator Loss: 6.6014024811266765, Generator Loss: 3.3405757545073285
Epoch 100/100, Val Disc Loss: 6.601365028865754, Val Gen Loss: 0.06155784108809062
Final Average L1 Loss: Test-: 6.159258015926868 Val-: 6.259872433951212 Train-: 3.6009003159055766
Testing DT - Last 10 Average L1 Loss Trend: Not Strictly Decreasing
Average SSIM Score for Training: 99.11197097605452
Average SSIM Score for Validation: 98.32972945494417
Average SSIM Score for testing: 98.29212238850143
```

Training input/Real images



Last Frames



Generated images/Fake images

